

Yarn 常见问题维护手册 V1.0

Yarn 常见问题维护手册 V1.0

文档版本 01
发布日期 2016-03-31

华为技术有限公司



版权所有 © 华为技术有限公司 2016。 保留一切权利。

非经本公司书面许可，任何单位和个人不得擅自摘抄、复制本文档内容的部分或全部，并不得以任何形式传播。

商标声明



和其他华为商标均为华为技术有限公司的商标。

本文档提及的其他所有商标或注册商标，由各自的所有人拥有。

注意

您购买的产品、服务或特性等应受华为公司商业合同和条款的约束，本文档中描述的全部或部分产品、服务或特性可能不在您的购买或使用范围之内。除非合同另有约定，华为公司对本文档内容不做任何明示或默示的声明或保证。

由于产品版本升级或其他原因，本文档内容会不定期进行更新。除非另有约定，本文档仅作为使用指导，本文档中的所有陈述、信息和建议不构成任何明示或暗示的担保。

华为技术有限公司

地址： 深圳市龙岗区坂田华为总部办公楼 邮编：518129

网址： <http://www.huawei.com>

客户服务邮箱： support@huawei.com

客户服务电话： 4008302118

目 录

| | |
|--|----|
| YARN | 2 |
| 1、基本概念 | 2 |
| 【概述】 | 2 |
| 【日志概述】 | 3 |
| 2、常见问题 | 7 |
| 【任务提交异常】 | 7 |
| [YARN-10001] 在 Windows 环境下提交 MR 任务失败。 | 7 |
| [YARN-10002] 怎样在 Yarn 提交任务的时候指定队列。 | 8 |
| [YARN-10003] 如何让 MR 任务输出的压缩文件。 | 8 |
| 【任务运行异常】 | 9 |
| [YARN-20001]MR 任务运行失败，报虚拟内存不足..... | 9 |
| [YARN-20002]MR 任务运行失败，报 OOM 异常..... | 10 |
| [YARN-20003] 集群资源足够时，仍然有大量任务处于 Accepted 状态 ... | 10 |
| [YARN-20004] OS 开启透明大页功能导致 MR 缓慢 | 12 |
| 【监控状态异常】 | 12 |
| [YARN-30001] NodeManager 启动时出现 Concerning。 | 12 |
| [Yarn-30002]在 OM 页面上添加 Yarn 的权限失败..... | 14 |
| [Yarn-30003]通过后台命令动态刷新队列配置失败 | 15 |
| [Yarn-30004]通过 FI Manager 管理页面动态刷新队列配置失败..... | 17 |
| 【咨询】 | 17 |
| [Yarn-40001]如何添加客户端参数..... | 17 |
| [Yarn-40002]如何添加服务端参数..... | 18 |

YARN

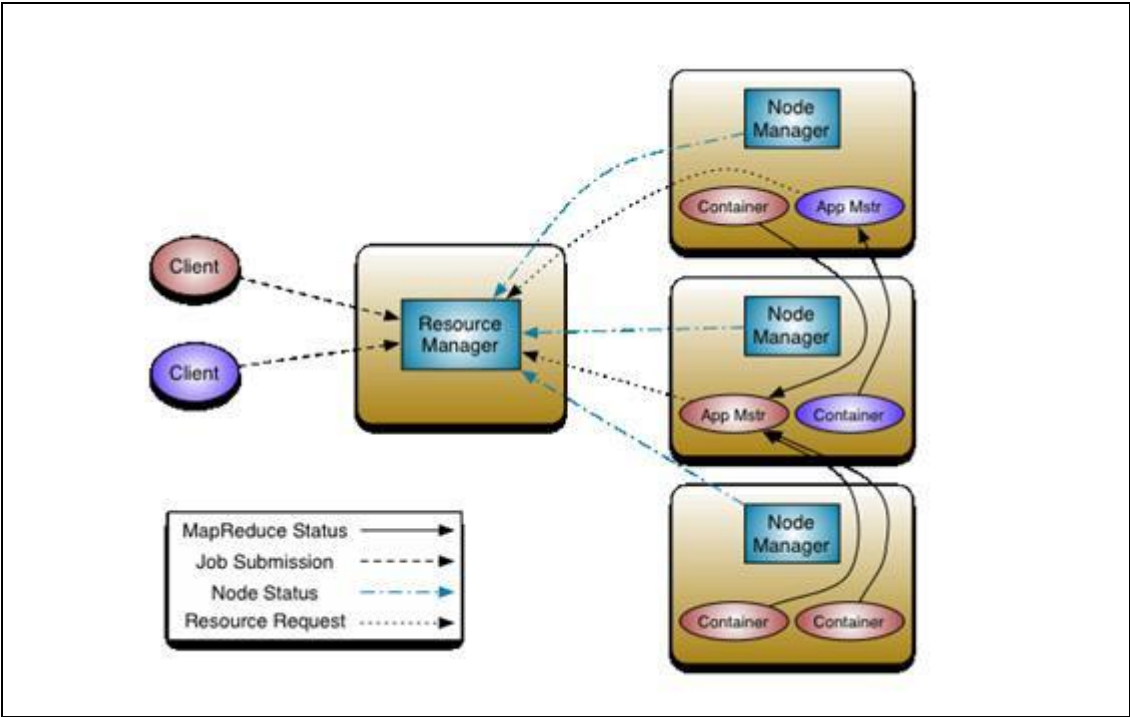
1、基本概念

本章简要介绍 Yarn 基本概念以及 Yarn 日志存放路径，以及常见问题的日志查看方法。

【概述】

Yarn 是 Hadoop 2.x 版本引入的一种资源管理框架。它将统一管理集群的资源，与任务调度分离，可以在其上层搭配不同的计算框架，如 MapReduce、Spark、Tez 等。

Yarn 结构图如下：



Yarn 的主要特点有：

- 将资源管理与任务调度分离，可以在其上层选择不同的分布式计算框架。
- 开放资源调度器的接口，可以选择使用不同的调度策略。

Yarn 各组件说明：

| | |
|-----------------|---------------------------------------|
| ResourceManager | Yarn 的主服务。主要完成资源申请、分配，任务的提交等调度任务。 |
| NodeManager | 以 container 的形式管理本节点上的资源，向 RM 保持心跳连接。 |

【日志概述】

Yarn 日志

Yarn 日志包括两类，运行日志、审计日志；

日志路径：

运行日志: `/var/log/Bigdata/yarn`

审计日志: `/var/log/Bigdata/audit/yarn`

各日志功能如下表：

| | | |
|------|---|------------------------------------|
| 运行日志 | yarn-<SSH_USER>- <process_name>- <hostname>.log | YARN 组件日志，记录 YARN 组件运行时候所产生的大部分日志。 |
| | yarn-<SSH_USER>- <process_name>- <hostname>.out | YARN 运行环境信息日志。 |
| | yarn.log | YARN 客户端操作日志。 |
| | prestart-Detail.log | 服务的预启动日志。 |
| | startDetail.log | 服务的启动日志。 |
| | stopDetail.log | 服务的停止日志。 |
| | hadoop.log | Hadoop 客户端操作日志。 |
| 审计日志 | yarn-audit- <process_name>.log | YARN 操作审计日志。 |

MapReduce 日志

日志路径：

MapReduce 任务在**运行的过程**中，Map 和 Reduce 任务所运行的 Container 日志位于该任务所在节点的以下路径：

`/srv/BigData/hadoop/data1/nm/containerlogs/<application_id>/<container_id>/`

| | | |
|--------------|--------|--------------------------------|
| Container 日志 | syslog | MapReduce 程序中使用 log4j 打印的日志信息。 |
| | stdout | MapReduce 程序中使用标准输出流打印的信息。 |
| | stderr | MapReduce 程序中使用标准错误流打印的信息。 |

MapReduce 任务**运行完成**后，相关日志会被拷贝至 HDFS 的以下路径，并且上述的三个子日志被合并为同一个文件：

`/tmp/logs/<user_name>/logs/<application_id>/<hostname_port>`

各个日志里都存放了什么信息？

审计日志

对于 Yarn 来说，审计日志主要记录了某个用户从哪个 IP 向 Yarn（主要是 ResourceManager）提交的任务信息，以及 AppMaster 向 ResourceManager 的注册信息。

如下的 RM 审计日志，表示在 2016-03-22 10:12:21 用户 admin 向 RM 提交了一个 MapReduce 任务，客户端 IP 为 xxx.xxx.1.18。其中主要包括了任务提交，任务的 APPMaster 注册，任务的结果三条记录。

```
2016-03-22 10:12:21,056 | INFO | IPC Server handler 13 on 26004 | USER=admin IP=xxx.xxx.1.18 OPERATION=Submit Application Request | TARGET=ClientRMService RESULT=SUCCESS APPID=application_1458609629121_0001 | RMAppAttemptImpl.java:91
2016-03-22 10:12:35,812 | INFO | Socket Reader #1 for port 26002 | Auth successful for appattempt_1458609629121_0001_000001 (auth:SIMPLE) | Server.java:1334
2016-03-22 10:12:35,885 | INFO | IPC Server handler 8 on 26002 | USER=admin IP=xxx.xxx.1.18 OPERATION=Register App Master | TARGET=ApplicationMasterService RESULT=SUCCESS APPID=application_1458609629121_0001 | RMAppAttemptImpl.java:91
2016-03-22 10:13:01,233 | INFO | Socket Reader #1 for port 26004 | Auth successful for hive/hadoop.hadoop.com@HADOOP.COM (auth:KERBEROS) | Server.java:1334
2016-03-22 10:13:05,722 | INFO | Socket Reader #1 for port 26004 | Auth successful for mapred/hadoop.hadoop.com@HADOOP.COM (auth:KERBEROS) | Server.java:1334
2016-03-22 10:13:12,236 | INFO | Socket Reader #1 for port 26004 | Auth successful for hive/hadoop.hadoop.com@HADOOP.COM (auth:KERBEROS) | Server.java:1334
2016-03-22 10:13:52,133 | INFO | AsyncDispatcher event handler | USER=admin OPERATION=Application Finished - Succeeded | TARGET=RMAppManager RESULT=SUCCESS APPID=application_1458609629121_0001 | RMAppAttemptImpl.java:91
```

审计日志就是对这些流程动作进行了一下记录，对 Yarn 的问题定位来说，较少来查看这些信息。

ResourceManager 运行日志

ResourceManager 负责接收客户端请求（MR/Spark 等任务），然后为任务分配第一个资源启动 APPMaster 并管理 APPMaster 的生命周期，接受来自 APPMaster 的资源申请。ResourceManager 运行日志中将记录这些信息。

以下为截取的 ResourceManager 日志，从内部状态机的角度展示了 RM 中 APP 的运行过程。

```
2016-03-22 11:09:14,059 | INFO | AsyncDispatcher event handler | application_1458609629121_0002 State change from NEW to NEW_SAVING | RMAppImpl.java:753
2016-03-22 11:09:14,072 | INFO | AsyncDispatcher event handler | application_1458609629121_0002 State change from NEW_SAVING to SUBMITTED | RMAppImpl.java:753
2016-03-22 11:09:14,074 | INFO | ResourceManager Event Processor | Accepted application application_1458609629121_0002 from user: admin, in queue: default | CapacityScheduler.java:772
2016-03-22 11:09:14,074 | INFO | AsyncDispatcher event handler | application_1458609629121_0002 State change from SUBMITTED to ACCEPTED | RMAppImpl.java:753
2016-03-22 11:09:14,075 | INFO | AsyncDispatcher event handler | appattempt_1458609629121_0002_000001 State change from NEW to SUBMITTED | RMAppAttemptImpl.java:800
2016-03-22 11:09:14,075 | INFO | ResourceManager Event Processor | Application application_1458609629121_0002 from user: admin activated in queue: default | LeafQueue.java:662
2016-03-22 11:09:14,075 | INFO | ResourceManager Event Processor | Added Application Attempt appattempt_1458609629121_0002_000001 to scheduler from user admin in queue default | CapacityScheduler.java:802
2016-03-22 11:09:14,078 | INFO | AsyncDispatcher event handler | appattempt_1458609629121_0002_000001 State change from SUBMITTED to SCHEDULED | RMAppAttemptImpl.java:800
2016-03-22 11:09:14,510 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000001 Container Transitioned from NEW to ALLOCATED | RMContainerImpl.java:394
2016-03-22 11:09:14,527 | INFO | AsyncDispatcher event handler | container_1458609629121_0002_01_000001 Container Transitioned from ALLOCATED to ACQUIRED | RMContainerImpl.java:394
2016-03-22 11:09:14,528 | INFO | AsyncDispatcher event handler | appattempt_1458609629121_0002_000001 State change from SCHEDULED to ALLOCATED_SAVING | RMAppAttemptImpl.java:800
2016-03-22 11:09:14,546 | INFO | AsyncDispatcher event handler | appattempt_1458609629121_0002_000001 State change from ALLOCATED_SAVING to ALLOCATED | RMAppAttemptImpl.java:800
2016-03-22 11:09:14,583 | INFO | AsyncDispatcher event handler | appattempt_1458609629121_0002_000001 State change from ALLOCATED to LAUNCHED | RMAppAttemptImpl.java:800
2016-03-22 11:09:15,514 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000001 Container Transitioned from ACQUIRED to RUNNING | RMContainerImpl.java:394
2016-03-22 11:09:36,342 | INFO | AsyncDispatcher event handler | appattempt_1458609629121_0002_000001 State change from LAUNCHED to RUNNING | RMAppAttemptImpl.java:800
2016-03-22 11:09:36,342 | INFO | AsyncDispatcher event handler | application_1458609629121_0002 State change from ACCEPTED to RUNNING | RMAppImpl.java:753
2016-03-22 11:09:37,702 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000002 Container Transitioned from NEW to ALLOCATED | RMContainerImpl.java:394
2016-03-22 11:09:37,712 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000003 Container Transitioned from NEW to RESERVED | RMContainerImpl.java:394
2016-03-22 11:09:38,028 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000004 Container Transitioned from NEW to ALLOCATED | RMContainerImpl.java:394
2016-03-22 11:09:38,030 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000005 Container Transitioned from NEW to ALLOCATED | RMContainerImpl.java:394
2016-03-22 11:09:38,032 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000006 Container Transitioned from NEW to RESERVED | RMContainerImpl.java:394
2016-03-22 11:09:38,052 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000007 Container Transitioned from NEW to ALLOCATED | RMContainerImpl.java:394
2016-03-22 11:09:38,055 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000008 Container Transitioned from NEW to ALLOCATED | RMContainerImpl.java:394
2016-03-22 11:09:38,496 | INFO | IPC Server handler 29 on 26002 | container_1458609629121_0002_01_000009 Container Transitioned from ALLOCATED to ACQUIRED | RMContainerImpl.java:394
2016-03-22 11:09:38,493 | INFO | IPC Server handler 29 on 26002 | container_1458609629121_0002_01_000004 Container Transitioned from ALLOCATED to ACQUIRED | RMContainerImpl.java:394
2016-03-22 11:09:38,496 | INFO | IPC Server handler 29 on 26002 | container_1458609629121_0002_01_000005 Container Transitioned from ALLOCATED to ACQUIRED | RMContainerImpl.java:394
2016-03-22 11:09:38,498 | INFO | IPC Server handler 29 on 26002 | container_1458609629121_0002_01_000006 Container Transitioned from ALLOCATED to ACQUIRED | RMContainerImpl.java:394
2016-03-22 11:09:38,498 | INFO | IPC Server handler 29 on 26002 | container_1458609629121_0002_01_000007 Container Transitioned from ALLOCATED to ACQUIRED | RMContainerImpl.java:394
2016-03-22 11:09:39,044 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000004 Container Transitioned from ACQUIRED to RUNNING | RMContainerImpl.java:394
2016-03-22 11:09:39,047 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000005 Container Transitioned from ACQUIRED to RUNNING | RMContainerImpl.java:394
2016-03-22 11:09:39,726 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000002 Container Transitioned from ACQUIRED to RUNNING | RMContainerImpl.java:394
2016-03-22 11:09:40,080 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000007 Container Transitioned from ACQUIRED to RUNNING | RMContainerImpl.java:394
2016-03-22 11:09:40,081 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000008 Container Transitioned from ACQUIRED to RUNNING | RMContainerImpl.java:394
2016-03-22 11:09:59,195 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000005 Container Transitioned from RUNNING to COMPLETED | RMContainerImpl.java:394
2016-03-22 11:09:59,201 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000004 Container Transitioned from RUNNING to COMPLETED | RMContainerImpl.java:394
2016-03-22 11:09:59,208 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000009 Container Transitioned from NEW to ALLOCATED | RMContainerImpl.java:394
2016-03-22 11:09:59,576 | INFO | IPC Server handler 27 on 26002 | container_1458609629121_0002_01_000008 Container Transitioned from ALLOCATED to ACQUIRED | RMContainerImpl.java:394
2016-03-22 11:10:03,251 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000009 Container Transitioned from ACQUIRED to RUNNING | RMContainerImpl.java:394
2016-03-22 11:10:09,974 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000002 Container Transitioned from RUNNING to COMPLETED | RMContainerImpl.java:394
2016-03-22 11:10:12,386 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000009 Container Transitioned from RUNNING to COMPLETED | RMContainerImpl.java:394
2016-03-22 11:10:12,386 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000005 Container Transitioned from RUNNING to COMPLETED | RMContainerImpl.java:394
2016-03-22 11:10:50,021 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000008 Container Transitioned from RUNNING to COMPLETED | RMContainerImpl.java:394
2016-03-22 11:11:01,063 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000010 Container Transitioned from RUNNING to COMPLETED | RMContainerImpl.java:394
2016-03-22 11:11:01,905 | INFO | AsyncDispatcher event handler | appattempt_1458609629121_0002_000001 State change from RUNNING to FINAL_SAVING | RMAppAttemptImpl.java:800
2016-03-22 11:11:01,922 | INFO | AsyncDispatcher event handler | application_1458609629121_0002 State change from RUNNING to FINAL_SAVING | RMAppImpl.java:753
2016-03-22 11:11:01,922 | INFO | AsyncDispatcher event handler | appattempt_1458609629121_0002_000001 State change from FINAL_SAVING to FINISHING | RMAppAttemptImpl.java:800
2016-03-22 11:11:01,971 | INFO | AsyncDispatcher event handler | application_1458609629121_0002 State change from FINAL_SAVING to FINISHING | RMAppImpl.java:753
2016-03-22 11:11:14,288 | INFO | ResourceManager Event Processor | container_1458609629121_0002_01_000001 Container Transitioned from RUNNING to COMPLETED | RMContainerImpl.java:394
2016-03-22 11:11:14,289 | INFO | AsyncDispatcher event handler | appattempt_1458609629121_0002_000001 State change from FINISHING to FINISHED | RMAppAttemptImpl.java:800
2016-03-22 11:11:14,289 | INFO | AsyncDispatcher event handler | application_1458609629121_0002 State change from FINISHING to FINISHED | RMAppImpl.java:753
```

NodeManager 运行日志

NodeManager 管理本节点的所有资源，向 RM 保持心跳并领取和启动对应的任务。NodeManager 运行日志将记录这些信息。

以下为截取的 NodeManager 日志，从内部状态机的角度展示了 RM 中 APP 的运行过程。

```
2016-03-22 11:17:57,293 | INFO | AsyncDispatcher event handler | Application application_1458609629121_0003 transitioned from NEW to INITING | ApplicationImpl.java:464
2016-03-22 11:17:57,432 | INFO | AsyncDispatcher event handler | Application application_1458609629121_0003 transitioned from INITING to RUNNING | ApplicationImpl.java:464
2016-03-22 11:17:57,433 | INFO | AsyncDispatcher event handler | Container container_1458609629121_0003_01_000001 transitioned from NEW to LOCALIZING | ContainerImpl.java:1129
2016-03-22 11:17:57,434 | INFO | AsyncDispatcher event handler | Resource hdfs://hacluster:8020/tmp/hadoop-yarn/staging/admin/.staging/job_1458609629121_0003/job.splitmetainfo transitioned from INIT to
DOWNLOADING | LocalizedResource.java:203
2016-03-22 11:17:57,434 | INFO | AsyncDispatcher event handler | Resource hdfs://hacluster:8020/tmp/hadoop-yarn/staging/admin/.staging/job_1458609629121_0003/job.jar transitioned from INIT to DOWNLOA
DS | LocalizedResource.java:203
2016-03-22 11:17:57,434 | INFO | AsyncDispatcher event handler | Resource hdfs://hacluster:8020/tmp/hadoop-yarn/staging/admin/.staging/job_1458609629121_0003/job.split transitioned from INIT to DOWNLOA
DS | LocalizedResource.java:203
2016-03-22 11:17:57,434 | INFO | AsyncDispatcher event handler | Resource hdfs://hacluster:8020/tmp/hadoop-yarn/staging/admin/.staging/job_1458609629121_0003/job.xml transitioned from INIT to DOWNLOA
DS | LocalizedResource.java:203
2016-03-22 11:18:03,597 | INFO | IPC Server handler 3 on 26007 | Resource hdfs://hacluster:8020/tmp/hadoop-yarn/staging/admin/.staging/job_1458609629121_0003/job.splitmetainfo(->srv/BigData/hadoop/d
atal/nm/localdir/usercache/admin/appcache/application_1458609629121_0003/filecache/10/job.splitmetainfo) transitioned from DOWNLOADING to LOCALIZED | LocalizedResource.java:203
2016-03-22 11:18:03,611 | INFO | IPC Server handler 4 on 26007 | Resource hdfs://hacluster:8020/tmp/hadoop-yarn/staging/admin/.staging/job_1458609629121_0003/job.jar(->srv/BigData/hadoop/datal/nm/loc
aldir/usercache/admin/appcache/application_1458609629121_0003/filecache/11/job.jar) transitioned from DOWNLOADING to LOCALIZED | LocalizedResource.java:203
2016-03-22 11:18:03,612 | INFO | IPC Server handler 1 on 26007 | Resource hdfs://hacluster:8020/tmp/hadoop-yarn/staging/admin/.staging/job_1458609629121_0003/job.split(->srv/BigData/hadoop/datal/nm/loc
aldir/usercache/admin/appcache/application_1458609629121_0003/filecache/12/job.split) transitioned from DOWNLOADING to LOCALIZED | LocalizedResource.java:203
2016-03-22 11:18:03,743 | INFO | IPC Server handler 0 on 26007 | Resource hdfs://hacluster:8020/tmp/hadoop-yarn/staging/admin/.staging/job_1458609629121_0003/job.xml(->srv/BigData/hadoop/datal/nm/loc
aldir/usercache/admin/appcache/application_1458609629121_0003/filecache/13/job.xml) transitioned from DOWNLOADING to LOCALIZED | LocalizedResource.java:203
2016-03-22 11:18:03,743 | INFO | AsyncDispatcher event handler | Container container_1458609629121_0003_01_000001 transitioned from LOCALIZING to LOCALIZED | ContainerImpl.java:1129
2016-03-22 11:18:03,812 | INFO | AsyncDispatcher event handler | Container container_1458609629121_0003_01_000001 transitioned from LOCALIZED to RUNNING | ContainerImpl.java:1129
2016-03-22 11:18:15,560 | INFO | AsyncDispatcher event handler | Container container_1458609629121_0003_01_000002 transitioned from NEW to LOCALIZING | ContainerImpl.java:1129
2016-03-22 11:18:15,560 | INFO | AsyncDispatcher event handler | Resource hdfs://hacluster:8020/tmp/hadoop-yarn/staging/admin/.staging/job_1458609629121_0003/job.jar transitioned from INIT to DOWNLOADING |
LocalizedResource.java:203
2016-03-22 11:18:15,560 | INFO | AsyncDispatcher event handler | Resource hdfs://hacluster:8020/tmp/hadoop-yarn/staging/admin/.staging/job_1458609629121_0003/job.xml transitioned from INIT to DOWNLOADING |
LocalizedResource.java:203
2016-03-22 11:18:15,560 | INFO | IPC Server handler 4 on 26007 | Resource hdfs://hacluster:8020/tmp/hadoop-yarn/staging/admin/.staging/job_1458609629121_0003/job.jar(->srv/BigData/hadoop/datal/nm/localdir
/usercache/admin/appcache/application_1458609629121_0003/filecache/14/job.jar) transitioned from DOWNLOADING to LOCALIZED | LocalizedResource.java:203
2016-03-22 11:18:15,560 | INFO | IPC Server handler 1 on 26007 | Resource hdfs://hacluster:8020/tmp/hadoop-yarn/staging/admin/.staging/job_1458609629121_0003/job.xml(->srv/BigData/hadoop/datal/nm/localdir
/usercache/admin/appcache/application_1458609629121_0003/filecache/15/job.xml) transitioned from DOWNLOADING to LOCALIZED | LocalizedResource.java:203
2016-03-22 11:18:25,284 | INFO | AsyncDispatcher event handler | Container container_1458609629121_0003_01_000002 transitioned from LOCALIZING to LOCALIZED | ContainerImpl.java:1129
2016-03-22 11:18:25,453 | INFO | AsyncDispatcher event handler | Container container_1458609629121_0003_01_000002 transitioned from LOCALIZED to RUNNING | ContainerImpl.java:1129
2016-03-22 11:18:39,121 | INFO | AsyncDispatcher event handler | Container container_1458609629121_0003_01_000002 transitioned from RUNNING to KILLING | ContainerImpl.java:1129
2016-03-22 11:18:39,134 | WARN | ContainerLauncher #17 | Exit code from container container_1458609629121_0003_01_000002 is : 143 | LinuxContainerExecutor.java:389
2016-03-22 11:18:39,952 | INFO | AsyncDispatcher event handler | Container container_1458609629121_0003_01_000002 transitioned from KILLING to CONTAINER_CLEANEDUP_AFTER_KILL | ContainerImpl.java:1129
2016-03-22 11:18:39,952 | INFO | AsyncDispatcher event handler | Container container_1458609629121_0003_01_000002 transitioned from CONTAINER_CLEANEDUP_AFTER_KILL to DONE | ContainerImpl.java:1129
2016-03-22 11:18:39,953 | INFO | AsyncDispatcher event handler | Removing container_1458609629121_0003_01_000002 from application application_1458609629121_0003 | ApplicationImpl.java:347
2016-03-22 11:18:42,147 | INFO | Node Status Updater | Removed completed containers from NM context: [container_1458609629121_0003_01_000002] | NodeStatusUpdaterImpl.java:529
2016-03-22 11:19:25,303 | INFO | AsyncDispatcher event handler | Container container_1458609629121_0003_01_000001 transitioned from RUNNING to EXITED_WITH_SUCCESS | ContainerImpl.java:1129
2016-03-22 11:19:25,339 | INFO | AsyncDispatcher event handler | Container container_1458609629121_0003_01_000001 transitioned from EXITED_WITH_SUCCESS to DONE | ContainerImpl.java:1129
2016-03-22 11:19:25,340 | INFO | AsyncDispatcher event handler | Removing container_1458609629121_0003_01_000001 from application application_1458609629121_0003 | ApplicationImpl.java:347
2016-03-22 11:19:26,273 | INFO | Node Status Updater | Removed completed containers from NM context: [container_1458609629121_0003_01_000001] | NodeStatusUpdaterImpl.java:529
2016-03-22 11:19:26,273 | INFO | AsyncDispatcher event handler | Application application_1458609629121_0003 transitioned from RUNNING to APPLICATION_RESOURCES_CLEANINGUP | ApplicationImpl.java:464
2016-03-22 11:19:26,274 | INFO | AsyncDispatcher event handler | Application application_1458609629121_0003 transitioned from APPLICATION_RESOURCES_CLEANINGUP to FINISHED | ApplicationImpl.java:464
```

GC 日志

ResourceManager 和 NodeManager 均有 GC 日志，当我们遇到 GC 问题可以 GC 日志以快速定位是否是 GC 导致。

MapReduce 日志

进入 OM 页面 Services->Yarn->ResourceManager,如图:

The screenshot shows the Yarn Summary page in the ResourceManager WebUI. The 'Health Status' is 'Good' and the 'Configuration Status' is 'Synchronized'. The 'Version' is '2.7.1'. There are '0' running and '0' pending applications. At the bottom, the 'ResourceManager(Active)' link is highlighted with a red box.

| Health Status | Configuration Status | Version | Running App | Pending App |
|---------------|----------------------|---------|-------------|-------------|
| Good | Synchronized | 2.7.1 | 0 | 0 |

ResourceManager(Standby)
ResourceManager(Active)

进入 Yarn 的页面后，找到我们刚才找到的 jobid，如下：

▼ Cluster

[About](#)
[Nodes](#)
[Applications](#)

[NEW](#)
[NEW SAVING](#)
[SUBMITTED](#)
[ACCEPTED](#)
[RUNNING](#)
[FINISHED](#)
[FAILED](#)
[KILLED](#)

[Scheduler](#)

► Tools

Cluster Metrics

| Apps Submitted | Apps Pending | Apps Running | Apps Completed | Containers Running | Memory Used | Memory Total | Memory Reserved |
|----------------|--------------|--------------|----------------|--------------------|-------------|--------------|-----------------|
| 1 | 0 | 0 | 1 | 0 | 0 B | 32 GB | 0 B |

Show 20 ▼ entries

| ID ▼ | User ↕ | Name ↕ | Application Type ↕ |
|--|--------|--|--------------------|
| application 1458616562037 0001 | admin | select count(*) from lineitem(Stage-1) | MAPREDUCE |
| application 1458025151195 0081 | admin | select count(*) from lineitem(Stage-1) | MAPREDUCE |
| application 1458025151195 0080 | admin | select count(*) from orders(Stage-1) | MAPREDUCE |
| application 1458025151195 0079 | admin | select count(*) from lineitem(Stage-1) | MAPREDUCE |
| application 1458025151195 0078 | admin | select count(*) from lineitem(Stage-1) | MAPREDUCE |

然后可以看见如下页面，点击 History。

| | |
|-------------------|--|
| User: | admin |
| Name: | select count(*) from lineitem(Stage-1) |
| Application Type: | MAPREDUCE |
| Application Tags: | |
| State: | FINISHED |
| FinalStatus: | SUCCEEDED |
| Started: | Fri Mar 18 09:20:43 +0800 2016 |
| Elapsed: | 2mins, 38sec |
| Tracking URL: | History |
| Diagnostics: | |

进入下面页面，可以点击打开 AM 日志，以及任意一个 map 或者 reduce 的日志。

| | | | | |
|--|------------------------------|--|-----------|------------|
| Job Name: select count(*) from lineitem(Stage-1) | | | | Job Over |
| User Name: admin | | | | |
| Queue: default | | | | |
| State: SUCCEEDED | | | | |
| Uberized: false | | | | |
| Submitted: Fri Mar 18 09:20:43 CST 2016 | | | | |
| Started: Fri Mar 18 09:20:53 CST 2016 | | | | |
| Finished: Fri Mar 18 09:23:21 CST 2016 | | | | |
| Elapsed: 2mins, 28sec | | | | |
| Diagnostics: | | | | |
| Average Map Time 17sec | | | | |
| Average Shuffle Time 1mins, 31sec | | | | |
| Average Merge Time 0sec | | | | |
| Average Reduce Time 1sec | | | | |
| ApplicationMaster | | | | |
| Attempt Number | Start Time | | Node | Logs |
| | Fri Mar 18 09:20:47 CST 2016 | | vm3:26010 | logs |
| Task Type | Total | | Complete | |
| Map | 29 | | 29 | |
| Reduce | 1 | | 1 | |
| Attempt Type | Failed | | Killed | Successful |
| Maps | 0 | | 0 | 29 |
| Reduces | 0 | | 0 | 1 |

最终可以看到 MapReduce 的执行日志：


```
Log Type: stderr
Log Length: 222
log4j:WARN No appenders could be found for logger (org.apache.hadoop.ipc.Server).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.

Log Type: stdout
Log Length: 0

Log Type: syslog
Log Length: 3972205
Showing 4096 bytes of 3972205 total. Click here for the full log.
DEBUG [IPC Server handler 8 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 8 on 25102: exiting
2016-03-18 09:23:22,627 DEBUG [IPC Server handler 12 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 12 on 25102: exiting
2016-03-18 09:23:22,627 DEBUG [IPC Server handler 13 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 13 on 25102: exiting
2016-03-18 09:23:22,627 DEBUG [IPC Server handler 14 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 14 on 25102: exiting
2016-03-18 09:23:22,627 DEBUG [IPC Server handler 15 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 15 on 25102: exiting
2016-03-18 09:23:22,627 DEBUG [IPC Server handler 16 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 16 on 25102: exiting
2016-03-18 09:23:22,627 DEBUG [IPC Server handler 17 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 17 on 25102: exiting
2016-03-18 09:23:22,627 DEBUG [IPC Server handler 18 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 18 on 25102: exiting
2016-03-18 09:23:22,627 DEBUG [IPC Server handler 19 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 19 on 25102: exiting
2016-03-18 09:23:22,627 DEBUG [IPC Server handler 20 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 20 on 25102: exiting
2016-03-18 09:23:22,627 DEBUG [IPC Server handler 21 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 21 on 25102: exiting
2016-03-18 09:23:22,628 DEBUG [IPC Server handler 23 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 23 on 25102: exiting
2016-03-18 09:23:22,628 DEBUG [IPC Server handler 22 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 22 on 25102: exiting
2016-03-18 09:23:22,628 DEBUG [IPC Server handler 24 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 24 on 25102: exiting
2016-03-18 09:23:22,629 DEBUG [IPC Server handler 27 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 27 on 25102: exiting
2016-03-18 09:23:22,629 DEBUG [IPC Server handler 25 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 25 on 25102: exiting
2016-03-18 09:23:22,629 INFO [IPC Server listener on 25102] org.apache.hadoop.ipc.Server: Stopping IPC Server listener on 25102
2016-03-18 09:23:22,629 DEBUG [IPC Server handler 26 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 26 on 25102: exiting
2016-03-18 09:23:22,629 DEBUG [IPC Server handler 28 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 28 on 25102: exiting
2016-03-18 09:23:22,629 DEBUG [IPC Server handler 29 on 25102] org.apache.hadoop.ipc.Server: IPC Server handler 29 on 25102: exiting
2016-03-18 09:23:22,633 DEBUG [IPC Server Responder] org.apache.hadoop.ipc.Server: Checking for old call responses.
2016-03-18 09:23:22,633 INFO [IPC Server Responder] org.apache.hadoop.ipc.Server: Stopping IPC Server Responder
2016-03-18 09:23:22,633 DEBUG [Thread-142] org.apache.hadoop.service.CompositeService: org.apache.hadoop.mapred.TaskAttemptListenerImpl: s
2016-03-18 09:23:22,633 DEBUG [Thread-142] org.apache.hadoop.service.CompositeService: Stopping service #0: Service TaskHeartbeatHandler i
```

2、常见问题

注意：案例中 IP 地址信息，请根据实际情况进行修改。

【任务提交异常】

[YARN-10001] 在 Windows 环境下提交 MR 任务失败。

【问题背景与现象】

在 windows 上开发 MR 应用程序后失败：

```
40014 [main] INFO org.apache.hadoop.mapreduce.Job - Job job_1458609629121_0009 failed with state FAILED
For more detailed output, check the application tracking page:https://160-138-1-182:26001/cluster/app/app
Diagnostics: Exception from container-launch.
Container id: container_1458609629121_0009_02_000001
Exit code: 1
Stack trace: ExitCodeException exitCode=1:
    at org.apache.hadoop.util.Shell.runCommand(Shell.java:556)
    at org.apache.hadoop.util.Shell.run(Shell.java:467)
    at org.apache.hadoop.util.Shell$ShellCommandExecutor.execute(Shell.java:733)
    at org.apache.hadoop.yarn.server.nodemanager.LinuxContainerExecutor.launchContainer(LinuxContainerExecutor.java:100)
    at org.apache.hadoop.yarn.server.nodemanager.containermanager.launcher.ContainerLaunch.call(ContainerLaunch.java:100)
    at org.apache.hadoop.yarn.server.nodemanager.containermanager.launcher.ContainerLaunch.call(ContainerLaunch.java:100)
    at java.util.concurrent.FutureTask.run(FutureTask.java:266)
    at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1142)
    at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:617)
    at java.lang.Thread.run(Thread.java:745)

Shell output: main : command provided 1
main : run as user is omm
main : requested yarn user is omm

|
Container exited with a non-zero exit code 1
Failing this attempt. Failing the application.
40075 [main] INFO org.apache.hadoop.mapreduce.Job - Counters: 0
```

【可能原因】

1. 未设置跨平台提交任务的配置参数 `mapreduce.app-submission.cross-platform`。

【解决办法】

1. 将客户端参数 `mapreduce.app-submission.cross-platform` 参数设置为 `true`。
参数添加方式请参阅[\[Yarn-40001\]](#)如何添加客户端参数。

[YARN-10002] 怎样在 Yarn 提交任务的时候指定队列。

【问题背景与现象】

如何在 Yarn 提交任务的时候指定队列？

【解决办法】

1. 将 `mapreduce.job.queueName` 设置为想要提交的队列名。
参数添加方式请参阅[\[Yarn-40001\]](#)如何添加客户端参数。

[YARN-10003] 如何让 MR 任务输出的压缩文件。

【问题背景与现象】

如何让 MR 任务输出的压缩文件？

【解决办法】

1. 当前 Yarn 支持以下几种压缩格式：

```
org.apache.hadoop.io.compress.BZip2Codec
org.apache.hadoop.io.compress.Lz4Codec
org.apache.hadoop.io.compress.DeflateCodec
org.apache.hadoop.io.compress.SnappyCodec
org.apache.hadoop.io.compress.GzipCodec
```

2. 以 gz 为例：通过在客户端代码或者 mapred-sit.xml 中指定如下参数可以让 MR 任务输出为 gz 压缩：

```
mapreduce.output.fileoutputformat.compress=true
mapreduce.output.fileoutputformat.compress.type=BLOCK
mapreduce.output.fileoutputformat.compress.codec=org.apache.hadoop.io.compress.GzipCodec
```

3. 可以仅配置 map 的输出是否为压缩文件

```
mapreduce.map.output.compress=true
mapreduce.map.output.compress.codec=org.apache.hadoop.io.compress.GzipCodec
```

4. 参数配置方式请参阅[Yarn-40001]如何添加客户端参数。

【任务运行异常】

[YARN-20001]MR 任务运行失败，报虚拟内存不足

【问题背景与现象】

MR 任务执行失败。

【原因分析】

1. 查看 AM 日志中发现 **beyond virtual memory limits**，具体信息如下：

```
Container[pid=41884,containerID=container_1405950053048_0016_01_000284] is
running beyond virtual memory limits. Current usage: 314.6 MB of 2.9 GB physical
memory used; 8.7 GB of 6.2 GB virtual memory used. Killing container.
```

2. 分配给 container 的虚拟内存不足

【解决办法】

1. 增加 container 能使用的虚拟内存大小。
虚拟内存的配置是由物理内存以及虚拟内存和物理内存的使用比例得到的 $\text{yarn.nodemanager.resource.memory-mb} * \text{yarn.nodemanager.vmem-pmem-ratio}$ ，因此，可以通过增加这两个参数的方式进行处理。其中 $\text{yarn.nodemanager.vmem-pmem-ratio}$ 是服务端参数。

【扩展介绍】

1. **MR 的内存参数介绍**

| | |
|--------------------------------------|---|
| yarn.app.mapreduce.am.resource.mb | MR ApplicationMaster 占用的内存量 |
| yarn.app.mapreduce.am.command-opts | AM 进程的 JVM 参数，其中-Xmx 最大不能超过 yarn.app.mapreduce.am.resource.mb |
| mapreduce.map.memory.mb | 每个 Map Task 需要的内存量 |
| mapreduce.map.java.opts | Map 的 JVM 参数，其中-Xmx 最多不能超过 mapreduce.map.memory.mb * 0.75 |
| mapreduce.reduce.memory.mb | 每个 Reduce Task 需要的内存量 |
| mapreduce.reduce.java.opts | Reduce 的 JVM 参数，其中-Xmx 最多不能超过 mapreduce.reduce.memory.mb * 0.75 |
| yarn.nodemanager.vmem-pmem-ratio | Container 能使用的虚拟内存和物理内存的比例 |
| yarn.scheduler.maximum-allocation-mb | 单个 Container 能申请的最大内存 |
| yarn.scheduler.minimum-allocation-mb | 单个 Container 能申请的最小内存 |

[YARN-20002]MR 任务运行失败，报 OOM 异常

【问题背景与现象】

MR 任务执行失败。

【原因分析】

1. 查看 AM 日志中发现 **beyond virtual memory limits**，具体信息如下：

```
Container[pid=41884,containerID=container_1405950053048_0016_01_000284] is
running beyond virtual memory limits. Current usage: 314.6 MB of 2.9 GB physical
memory used; 8.7 GB of 6.2 GB virtual memory used. Killing container.
```

2. 分配给 container 的内存不足

【解决办法】

1. 增加 container 能使用的内存大小。
增加 mapreduce.map.memory.mb 或 mapreduce.map.java.opts 的值。这两者是有关联关系的，通常默认是 0.8，建议设置为 0.75，可由 mapreduce.job.heap.memory-mb.ratio 进行指定。

[YARN-20003] 集群资源足够时，仍然有大量任务处于 Accepted 状态

【问题背景与现象】

集群资源足够时，仍然有大量任务处于 Accepted 状态。

Cluster

About
Nodes
Node Labels
Applications
NEW
NEW SAVING
SUBMITTED
ACCEPTED
RUNNING
FINISHED
FAILED
KILLED
Scheduler
Tools

Cluster Metrics

| Apps Submitted | Apps Pending | Apps Running | Apps Completed | Containers Running | Memory Used | Memory Total | Memory Reserved | VCores Used | VCores Total | VCores Reser |
|----------------|--------------|--------------|----------------|--------------------|-------------|--------------|-----------------|-------------|--------------|--------------|
| 727 | 6 | 3 | 718 | 5 | 7.50 GB | 40 GB | 0 B | 5 | 40 | 0 |

Scheduler Metrics

| Scheduler Type | Scheduling Resource Type | Minimum Allocated |
|--------------------|--------------------------|------------------------|
| Capacity Scheduler | [MEMORY, CPU] | <memory:512, vCores:1> |

Show 20 entries

| ID | User | Name | Application Type | Queue | StartTime | FinishTime | State | FinalStatus |
|--------------------------------|-------|--|------------------|---------|--------------------------------|--------------------------------|----------|-------------|
| application 1447638446057_0745 | admin | Collect Female Info | MAPREDUCE | default | Tue Nov 17 20:47:18 +0800 2015 | N/A | ACCEPTED | UNDEFINED |
| application 1447638446057_0744 | admin | Collect Female Info | MAPREDUCE | default | Tue Nov 17 20:47:13 +0800 2015 | N/A | ACCEPTED | UNDEFINED |
| application 1447638446057_0743 | admin | Collect Female Info | MAPREDUCE | default | Tue Nov 17 20:47:07 +0800 2015 | N/A | ACCEPTED | UNDEFINED |
| application 1447638446057_0742 | admin | Collect Female Info | MAPREDUCE | default | Tue Nov 17 20:47:02 +0800 2015 | N/A | ACCEPTED | UNDEFINED |
| application 1447638446057_0741 | admin | Collect Female Info | MAPREDUCE | default | Tue Nov 17 20:46:57 +0800 2015 | N/A | ACCEPTED | UNDEFINED |
| application 1447638446057_0740 | admin | Collect Female Info | MAPREDUCE | default | Tue Nov 17 20:46:53 +0800 2015 | N/A | ACCEPTED | UNDEFINED |
| application 1447638446057_0739 | admin | Collect Female Info | MAPREDUCE | default | Tue Nov 17 20:46:48 +0800 2015 | Tue Nov 17 20:47:27 +0800 2015 | FINISHED | SUCCEEDED |
| application 1447638446057_0738 | admin | select length(load_date)...length(load_date) (Stage-1) | MAPREDUCE | default | Tue Nov 17 20:46:43 +0800 2015 | Tue Nov 17 20:47:28 +0800 2015 | FINISHED | SUCCEEDED |

【原因分析】

1. 查看该队列的详细信息，查看 AM 最大允许资源配置是否已经超过。

Cluster
About
Nodes
Node Labels
Applications
NEW
NEW SAVING
SUBMITTED
ACCEPTED
RUNNING
FINISHED
FAILED
KILLED
Scheduler
Tools

Cluster Metrics

| | | | | | | | | | | |
|----------------|--------------|--------------|----------------|--------------------|-------------|--------------|-----------------|-------------|--------------|-----------------|
| Apps Submitted | Apps Pending | Apps Running | Apps Completed | Containers Running | Memory Used | Memory Total | Memory Reserved | VCores Used | VCores Total | VCores Reserved |
| 1 | 0 | 1 | 0 | 3 | 4.50 GB | 24 GB | 0 B | 3 | 48 | 0 |

Scheduler Metrics

| | | |
|--------------------|--------------------------|------------------------|
| Scheduler Type | Scheduling Resource Type | Minimum Allocation |
| Capacity Scheduler | [MEMORY, CPU] | <memory:512, vCores:1> |

Dump scheduler logs 1 min
Application Queues
Legend: Capacity Used Used (over capacity) Max Capacity

Queue: root
QueueA
QueueB
QueueC
root-default
xxx
default

Queue State: RUNNING
Used Capacity: 93.8%
Absolute Used Capacity: 18.8%
Absolute Capacity: 20.0%
Absolute Max Capacity: 100.0%
Used Resources: <memory:4608, vCores:3>
Num Schedulable Applications: 1
Num Non-Schedulable Applications: 0
Num Containers: 3
Max Applications: 2000
Max Applications Per User: 20000
Max Application Master Resources: <memory:2560, vCores:5>
Used Application Master Resources: <memory:1536, vCores:1>
Max Application Master Resources Per User: <memory:5120, vCores:10>

【解决办法】

1. 如果 Used Application Master Resources 已经超过 Max Application Master Resources,

则需要等待某个任务结束后才能继续启动下一个任务。可以调整 yarn.scheduler.capacity.maximum-am-resource-percent 参数，最大允许 AM 占用集群总资源的比例。

[YARN-20004] OS 开启透明大页功能导致 MR 缓慢

【问题背景与现象】

Map/reduce 任务时间特别长，并且 top 命令发现 cpu 占用率的 sy 项特别高（红框值超过 60%以上），OS 运行缓慢。

```
top - 11:13:41 up 29 days, 23:49, 3 users, load average: 8.81, 10.21, 10.84
Tasks: 268 total, 3 running, 264 sleeping, 0 stopped, 1 zombie
Cpu0  : 65.1%us, 18.7%sy, 0.0%ni, 0.0%id, 0.6%wa, 0.0%hi, 0.0%si, 15.7%st
Cpu1  : 76.6%us, 22.0%sy, 0.7%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.7%st
Cpu2  : 59.4%us, 40.6%sy, 0.0%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu3  : 67.8%us, 29.4%sy, 0.0%ni, 0.7%id, 0.0%wa, 0.7%hi, 0.0%si, 1.4%st
Mem: 17366728k total, 16173416k used, 1193312k free, 281500k buffers
Swap: 0k total, 0k used, 0k free, 2241096k cached

  PID USER      PR  NI  VIRT  RES  SHR  S %CPU  %MEM    TIME+  COMMAND
 13130 omm        20   0 2740m 152m 19m  S  90.0   0.9   0:05.50 java
 14122 omm        20   0 6742m  46m 20m  S  65.8   0.3   0:00.78 java
 14208 omm        20   0 6768m  30m 11m  S  65.8   0.2   0:00.64 java
 31718 omm        20   0 6066m 1.5g 19m  S   6.6   9.4  1105:57 java
 14426 root        20   0  104m 2396 1072  S   4.4   0.0   0:00.04 sudoExecute.sh
 18387 omm        20   0 5921m 358m 138m  S   4.4   2.1   3:57.45 java
 23076 omm        20   0 3097m 583m 19m  S   4.4   3.4  142:59.80 java
```

【原因分析】

1. OS 开启透明大页功能。

【解决办法】

1. 执行以下命令关闭透明大页功能：

```
vi /etc/rc.d/rc.local
```

```
echo "never" > /sys/kernel/mm/redhat_transparent_hugepage/enabled
```

```
echo "never" > /sys/kernel/mm/redhat_transparent_hugepage/defrag
```

```
echo "no" > /sys/kernel/mm/redhat_transparent_hugepage/khugepaged/defrag
```

【监控状态异常】

[YARN-30001] NodeManager 启动时出现 Concerning。

【问题背景与现象】

启动 Yarn 后，发现个别 NodeManager 出现 Concerning 状态。

【可能原因】

1. NodeManager 启动内存过小
2. NodeManager 依赖文件权限错误
3. NodeManager 读写磁盘错误

【原因分析】

1. 参看 NodeManager 的运行日志，出现 `doesn't satisfy minimum allocations` 错误，具体内容如下：

```
2016-03-03 16:53:35,102 | WARN | main | USER=mapred OPERATION=nmStartup  
TARGET=NodeManager RESULT=FAILURE DESCRIPTION=Exception occurred during  
startup | NMAuditLogger.java:288  
2016-03-03 16:53:35,102 | FATAL | main | Error starting NodeManager |  
NodeManager.java:552  
org.apache.hadoop.yarn.exceptions.YarnRuntimeException:  
org.apache.hadoop.yarn.exceptions.YarnRuntimeException: Recieved SHUTDOWN signal  
from Resourcemanager, Registration of NodeManager failed, Message from  
ResourceManager: NodeManager from xxx-xxx-1-182 doesn't satisfy minimum  
allocations, Sending SHUTDOWN signal to the NodeManager.
```

说明设置的 NodeManager 启动内存过小，小于 NodeManager 所需的最小内存解决方法参考 [NODEMANAGER 内存过小](#)

2. 检查

`opt/huawei/Bigdata/FusionInsight-Hadoop-2.7.1/hadoop/bin/container-executor` 权限，发现被修改

```
total 268  
--Sr-s---. 1 root wheel 114035 Sep 30 00:28 container-executor  
-rw-----. 1 root root 2 Feb 27 16:34 drop_caches  
-r-x-----. 1 omm wheel 6524 Sep 30 00:28 hadoop  
-r-x-----. 1 omm wheel 8514 Sep 30 00:28 hadoop.cmd  
-rwx-----. 1 omm wheel 2375 Oct 16 03:19 hadoop-oam.sh  
-r-x-----. 1 omm wheel 12523 Sep 30 00:28 hdfs  
-rwx-----. 1 omm wheel 1133 Oct 16 03:19 hdfs-backup.sh
```

需要将权限修改正确，解决方法参考 [NODEMANAGER 文件权限错误](#)

3. 查看 `/etc/fstab` 中是否包含该项目

```
#  
# /etc/fstab  
# Created by anaconda on Fri Apr 24 16:31:37 2015  
#  
# Accessible filesystems, by reference, are maintained under '/dev/disk'  
# See man pages fstab(5), findfs(8), mount(8) and/or blkid(8) for more info  
#  
/dev/mapper/vg_160138118-lv_root / ext4 defaults,noatime,errors=panic 1 1  
UUID=53befa2b-f7ac-43be-b013-3a1fbca267d2 /boot ext4 defaults,noatime,errors=panic 1 2  
/dev/mapper/vg_160138118-lv_home /home ext4 defaults,noatime,errors=panic 1 2  
/dev/mapper/vg_160138118-lv_swap swap swap defaults 0 0  
tmpfs /dev/shm tmpfs defaults 0 0  
devpts /dev/pts devpts gid=5,mode=620 0 0  
sysfs /sys sysfs defaults 0 0  
proc /proc proc defaults 0 0  
/dev/sdb1 /opt ext4 defaults,nosuid,noatime,errors=panic 0 0  
/dev/sdal /srv/BigData ext4 defaults,noatime 0 0  
~  
~
```


解决方法参考 [NODEMANAGER 磁盘挂载错误](#)

【解决办法】

1. NodeManager 内存过小

修改 `yarn.nodemanager.resource.memory-mb` 参数，至少需要大于 `yarn.scheduler.minimum-allocation-mb` 这个参数

2. NodeManager 文件权限错误

`chmod 050 container-executor`
`chmod +s container-executor`

3. NodeManager 磁盘挂载错误


确保停止该磁盘上所有文件的读写，然后删除红框中的项，执行 `mount -a` 命令重新挂载磁盘。

```
#  
# /etc/fstab  
# Created by anaconda on Fri Apr 24 16:31:37 2015  
#  
# Accessible filesystems, by reference, are maintained under '/dev/disk'  
# See man pages fstab(5), findfs(8), mount(8) and/or blkid(8) for more info  
#  
/dev/mapper/vg_160138118-lv_root / ext4 defaults,noatime,errors=panic 1 1  
UUID=53bafa2b-f7ac-43be-b013-3a1fbca267d2 /boot ext4 defaults,noatime,errors=panic 1 2  
/dev/mapper/vg_160138118-lv_home /home ext4 defaults,noatime,errors=panic 1 2  
/dev/mapper/vg_160138118-lv_swap swap swap defaults 0 0  
tmpfs /dev/shm tmpfs defaults 0 0  
devpts /dev/pts devpts gid=5,mode=620 0 0  
sysfs /sys sysfs defaults 0 0  
proc /proc proc defaults 0 0  
/dev/sdb1 /opt ext4 defaults nosuid,noatime,errors=panic 0 0  
/dev/sda1 /srv/BigData ext4 defaults,noatime 0 0  
~
```

[Yarn-30002]在 OM 页面上添加 Yarn 的权限失败

【问题背景与现象】

在 OM 页面创建角色时，添加 Yarn 权限失败：



【原因分析】

1. 查看/var/log/Bigdata/controller/aos.log，发现 **set role permission failed** 报，具体内容如下：

```
2016-03-23 11:11:40,877 INFO [qtp1768471788-1309 - /rolexml/role/addrole] Finish expand config items, expand type is MODIFY. com.huawei.security.aos.pluginmanager.pluginutil.restclientConfigItems(RestClient.java:1425)
2016-03-23 11:11:40,877 ERROR [qtp1768471788-1309 - /rolexml/role/addrole] Failed to set acl permissions for queue com.huawei.hadoop.om.security.aos.plugin.yarn.capacity.CaConfiguratio
roePermission(CaConfigurationExecutor.java:141)
2016-03-23 11:11:40,877 INFO [qtp1768471788-1309 - /rolexml/role/addrole] finish set service Yarn resource permission. com.huawei.security.aos.permmanager.PermManager.setResourcePerm
java:1080)
2016-03-23 11:11:40,877 ERROR [qtp1768471788-1309 - /rolexml/role/addrole] set service Yarn resource permission failed. com.huawei.security.aos.permmanager.PermManager.setResourcePerm
java:1085)
2016-03-23 11:11:40,878 ERROR [qtp1768471788-1309 - /rolexml/role/addrole] set role Yarn_test permission failed. com.huawei.security.aos.rolemanager.RoleManager.setRolePermissions(Rol
2016-03-23 11:11:40,878 ERROR [qtp1768471788-1309 - /rolexml/role/addrole] set role permission failed. com.huawei.security.aos.rolemanager.RoleManager.addRole(RoleManager.java:210)
2016-03-23 11:11:49,500 INFO [pool-1-thread-1] begin reload plugin file. com.huawei.security.aos.pluginmanager.pluginload.PluginClassLoaderReloadTask.run(PluginClassLoader.java:471)
2016-03-23 11:11:49,500 INFO [pool-1-thread-1] begin reload plugin file. com.huawei.security.aos.pluginmanager.pluginload.PluginClassLoaderReloadTask.run(PluginClassLoader.java:471)
```

文件权限错误导致添加用户失败。

【解决办法】

在 RM 主备节点上使用

1. `Isolf /opt/huawei/Bigdata/etc/*_*_ResourceManager/capacity-scheduler.xml` 查看是否有其他进程打开这个文件
2. `Isolf /opt/huawei/Bigdata/etc/*_*_ResourceManager/.capacity-scheduler.xml.swp` 查看是否有其他进程打开这个文件
3. 如果有，关闭这个进程。

[Yarn-30003]通过后台命令动态刷新队列配置失败

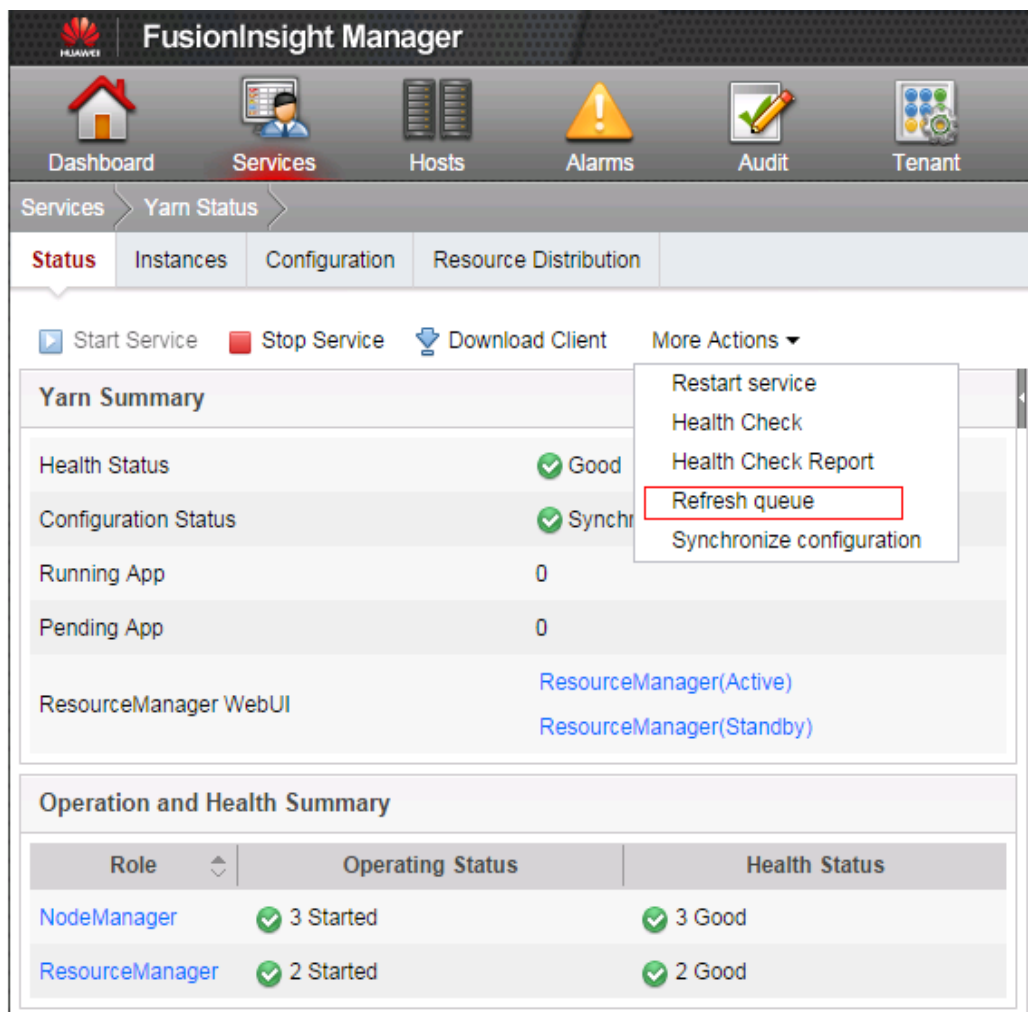
【问题背景与现象】

在 OM 页面创建角色时，添加 Yarn 权限失败：

```
sdp@dkfzxsmis0app57:~$ source /datafs2/sdp/FusionInsight_Client/bigdata_env
sdp@dkfzxsmis0app57:~$ kinit -k -t /datafs2/sdp/FusionInsight_Client/keytab/SDP_BATCH.keytab SDP_BATCH
sdp@dkfzxsmis0app57:~$ $HADOOP_CLIENT_NAME/HDFS/hadoop/bin/yarn rmadmin -refreshQueues
No GC PROFILE is given. Defaults to medium.
Exception in thread "main" java.lang.IllegalArgumentException: Can't get Kerberos realm
    at org.apache.hadoop.security.HadoopKerberosName.setConfiguration(HadoopKerberosName.java:65)
    at org.apache.hadoop.security.UserGroupInformation.initialize(UserGroupInformation.java:261)
    at org.apache.hadoop.security.UserGroupInformation.ensureInitialized(UserGroupInformation.java:246)
    at org.apache.hadoop.security.UserGroupInformation.loginUserFromSubject(UserGroupInformation.java:746)
    at org.apache.hadoop.security.UserGroupInformation.getLoginUser(UserGroupInformation.java:731)
    at org.apache.hadoop.security.UserGroupInformation.getCurrentUser(UserGroupInformation.java:633)
    at org.apache.hadoop.hdfs.tools.GetConf.run(GetConf.java:315)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:75)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:90)
    at org.apache.hadoop.hdfs.tools.GetConf.main(GetConf.java:332)
Caused by: java.lang.reflect.InvocationTargetException
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:57)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:606)
    at org.apache.hadoop.security.authentication.util.KerberosUtil.getDefaultRealm(KerberosUtil.java:75)
    at org.apache.hadoop.security.HadoopKerberosName.setConfiguration(HadoopKerberosName.java:63)
    ... 9 more
Caused by: KrbException: Cannot locate default realm
    at sun.security.krb5.Config.getDefaultRealm(Config.java:1183)
    ... 15 more
Caused by: KrbException: Generic error (description in e-text) (60) - Unable to locate Kerberos realm
    at sun.security.krb5.Config.getRealmFromDNS(Config.java:1279)
    at sun.security.krb5.Config.getDefaultRealm(Config.java:1164)
    ... 15 more
```

【解决办法】

FusionInsight 提供了图形化刷新队列的方案，不支持后台命令行刷新，可以使用如下操作刷新队列配置。



【扩展介绍】

Yarn 默认资源调度器主要参数介绍

| | |
|--|--|
| yarn.scheduler.capacity.<queue-path>.capacity | 队列的资源容量（百分比），所有队列的该配置加起来不能超过 100 |
| yarn.scheduler.capacity.<queue-path>.maximum-capacity | 队列的资源使用上限（百分比）。由于存在资源共享，因此一个队列使用的资源量可能超过其容量，而最多使用资源量可通过该参数限制。 |
| yarn.scheduler.capacity.<queue-path>.maximum-applications | 集群或者队列中同时处于等待和运行状态的应用程序数目上限，这是一个强限制，一旦集群中应用程序数目超过该上限，后续提交的应用程序将被拒绝，默认值为 10000 |
| yarn.scheduler.capacity.<queue-path>.maximum-am-resource-percent | 集群中用于运行应用程序 ApplicationMaster 的资源比例上限，该参数通常用于限制处于活动状态的应用程序数目。该参数类型为浮点型，默认是 0.1，表 |

| | |
|--|--------------------|
| | 示 10%。 |
| yarn.scheduler.capacity.<queue-path>.user-limit-factor | 每个用户最多可使用的资源量（百分比） |

[Yarn-30004]通过 FI Manager 管理页面动态刷新队列配置失败

【问题背景与现象】

FusionInsight Manager 管理页面刷新队列失败。

【原因分析】

RM 日志中打印 **Illegal capacity** 错误，红框处的值大于 1

```

s=0, numContainers=0 | CapacityScheduler.java:656
2016-03-26 10:56:22,546 | WARN | IPC Server handler 0 on 26005 | Exception refresh queues. | AdminService.java:710
java.io.IOException: Failed to re-init queues
    at org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.CapacityScheduler.reinitialize(CapacityScheduler.java:390)
    at org.apache.hadoop.yarn.server.resourcemanager.AdminService.refreshQueues(AdminService.java:376)
    at org.apache.hadoop.yarn.server.api.impl.pb.service.ResourceManagerAdministrationProtocolPBServiceImpl.refreshQueues(ResourceMan
    at org.apache.hadoop.yarn.proto.ResourceManagerAdministrationProtocol$ResourceManagerAdministrationProtocolService$2.callBlocking
    at org.apache.hadoop.ipc.ProtobufRpcEngine$Server$ProtoBufRpcInvoker.call(ProtobufRpcEngine.java:616)
    at org.apache.hadoop.ipc.RPC$Server.call(RPC.java:973)
    at org.apache.hadoop.ipc.Server$Handler$1.run(Server.java:2088)
    at org.apache.hadoop.ipc.Server$Handler$1.run(Server.java:2084)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1672)
    at org.apache.hadoop.ipc.Server$Handler.run(Server.java:2082)
Caused by: java.lang.IllegalArgumentException: Illegal capacity of 1.4000001 for children of queue root
    at org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.ParentQueue.setChildQueues(ParentQueue.java:145)
    at org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.CapacityScheduler.parseQueue(CapacityScheduler.java:646)
    at org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.CapacityScheduler.reinitializeQueues(CapacityScheduler.java:5
    at org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.CapacityScheduler.reinitialize(CapacityScheduler.java:386)
    ... 11 more
2016-03-26 10:56:22,554 | WARN | IPC Server handler 0 on 26005 | USER=mapred IP=160.138.1.181 OPERATION=refreshQueues TARGET=Ad
as. PERMISSIONS= | RMAuditLogger.java:51

```

【解决办法】

所有的 yarn.scheduler.capacity.<queue-path>.capacity 配置项的值加起来不得超过 100。

【咨询】

[Yarn-40001]如何添加客户端参数

【解决办法】

方式一：为客户端添加 mapred-site.xml 配置文件，在其中配置需要的属性，如下：

```

<property>
  <name>mapreduce.app-submission.cross-platform</name>
  <value>true</value>
</property>

```

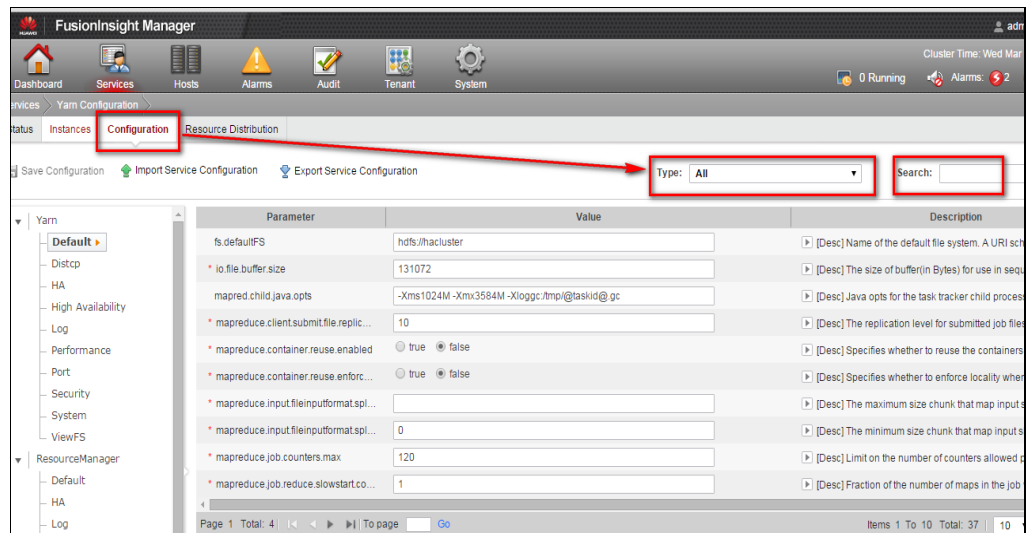
方式二：在客户端代码中添加参数设置代码，例如：

```
conf.set("mapreduce.app-submission.cross-platform", "true");
```

[Yarn-40002]如何添加服务端参数

【解决办法】

点击 Services->Yarn->Configuration，并选择 Type 为 ALL，如下：



在搜索框内搜索需要修改的参数并修改。